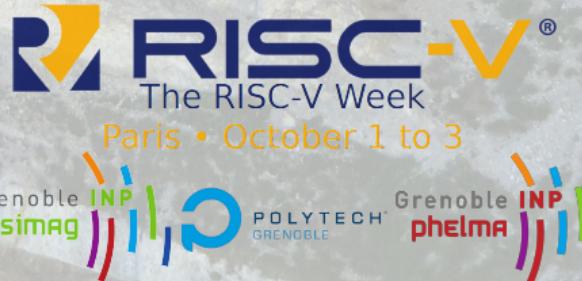


Teaching basic computer architecture, assembly language programming, and operating system design using RISC-V

Liliana Andrade, Mounir Benabdenbi,
Olivier Muller, Frédéric Rousseau,
Frédéric Pérot

tima.imag.fr/sls/people/petrot

frederic.petrot@univ-grenoble-alpes.fr



Outline

1 Preamble

2 Introduction

3 Overview of the curriculum

4 Class specifics

5 Take away

Preamble

Today's view on computers and processors

What computer a typical first year Computer Science (CS) students owns?

- ▶ You name it! But all are x86-64 based!

What processors do typical first year CS students know?

- ▶ Intel Core Ix, sometimes named x86 (or x86-64 for the connoisseurs)
- ▶ AMD Ryzen, but wait, this is also an x86!
- ▶ Atom, that's in my tablet, x86, uh, ...

Other devices

they know phone brands, but not the processors in them

1991

1992 1993 ... 2015 2016

2017

Intel did it right!



Preamble

Intel did it right, but ...

- ▶ "[x86] mov is Turing-complete", Stephen Dolan, 2013.
Actual code generator <https://github.com/xoreaxeaxeax/movfuscator>
But wait, why bother with instructions?
- ▶ "[x86] Page-faults are Turing-complete", Julian Bangert and Sergey Bratus, 2015.
Actual code generator <https://github.com/jbangert/trapcc>

Although x86 is the mainstream desktop computer architecture, it may be worth using something else as a pedagogical vehicle!

Outline

1 Preamble

2 Introduction

3 Overview of the curriculum

4 Class specifics

5 Take away

Introduction

Teaching computer engineering using RISC-V

Target students

Bologna style L3/M1/M2 in Computer Science or Electrical Engineering

In reality engineering schools in Grenoble

Ensimag, computer science

Phelma, electronic and micro-electronic engineering

Polytech Grenoble, electronics engineering and industrial IT

Goal

Unifying a set of classes using various processors under the RISC-V umbrella

Outline

- 1 Preamble
- 2 Introduction
- 3 Overview of the curriculum
- 4 Class specifics
- 5 Take away

Overview of the curriculum

Basic computer architecture

(Computer Science, 3rd year, ≈ 250 students, Electronics Engineering, 5th year, ≈ 20 students)

- ▶ digital circuit design for computer scientists
- ▶ ISA interpretation using a finite state machine + data-path

Assembly language programming

(Computer Science, 3rd year, ≈ 250 students)

- ▶ basic instruction usage
- ▶ function calling conventions and C ABI

Overview of the curriculum

Computer architecture

(Computer Science, 4th year, ≈ 40 students)

- ▶ 5 pipeline stages processor
- ▶ multiprocessor and atomic operations

Operating system implementation

(Computer Science, 4th year, ≈ 75 students)

- ▶ boot, interrupt, kernel threads
- ▶ virtual memory, processes

Overview of the curriculum

System level design in SystemC
(Computer Science, 5th year, \approx 20 students)

- ▶ hardware modeling in SystemC, transaction level modeling
- ▶ modeling in SystemC with native or cross-compiled software

HW/SW system integration
(Electrical Engineering, 5th year, \approx 40 students)

- ▶ performance analysis of a cache-based multiprocesssing system
- ▶ HW/SW integration on FPGA

Overview of the curriculum

	L3/Bachelor	M1/Master	M2/Master
Ensimag	Basic computer architecture	Operating system implementation	System level design in SystemC
	Assembly language programming	Computer architecture	
Polytech		Computer architecture	Basic computer architecture
Phelma			HW/SW system integration

Outline

- 1 Preamble
- 2 Introduction
- 3 Overview of the curriculum
- 4 Class specifics
- 5 Take away

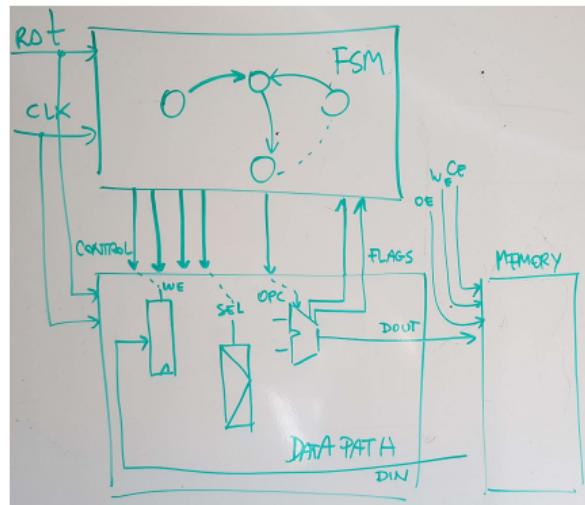
Basic computer architecture

Class and practical objectives

Write VHDL code interpreting instructions to understand why a computer needs to be powered, how it might execute a program, and why it is not indefinitely fast

What do we provide the students?

- ▶ VHDL of an FSM skeleton and a data-path with a few missing parts
- ▶ ordered list of instructions (encoding+behavior) to implement
- ▶ test environment to check their implementation



Basic computer architecture

What do we expect from the students?

- ▶ Proper decoding and execution of instructions with FSM+data-path
- ▶ Add missing parts in the data-path
 - ▶ condition computations for branches
 - ▶ control and status registers
 - ▶ interruptions
- ▶ Mapping on Xilinx FPGA board with 100 MHz minimal frequency target
- ▶ Own unitary test for each instruction
 - ▶ shall not depend upon future implemented instructions
 - ▶ are pretty easy develop as the micro-architecture is naïve

Basic computer architecture

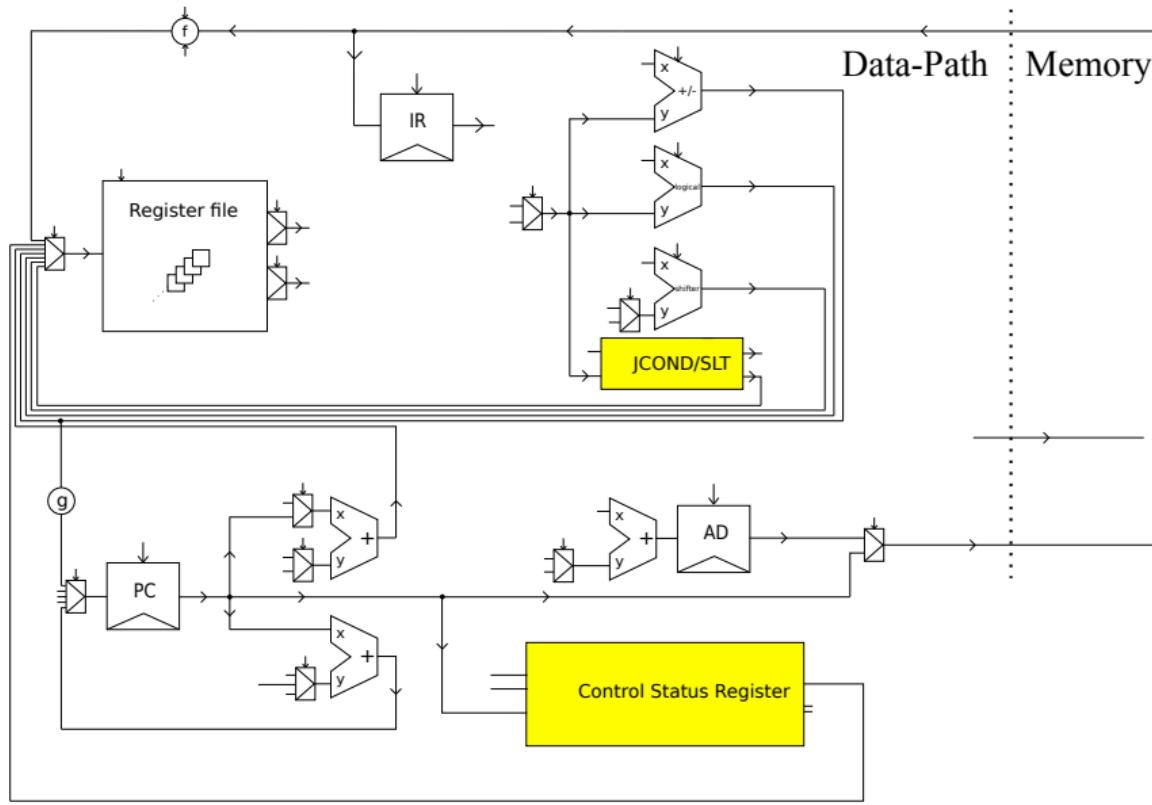
MIPS I neat, not so hype anymore

- ▶ all instructions are 32-bit
- ▶ 38 instructions including eret
- ▶ 3 instruction formats
- ▶ 5 immediate formats, zero or sign extended, 16-bit or 26-bit
- ▶ reg/reg instructions update rd,
reg/imm instructions update rt
⇒ sllv and srav instead of slli/srai
- ▶ delay slot, hidden from the students
- ▶ branch target computed using pc + 4
- ⇒ both former specificities assume pipeline implementations

RISC-V rv32i stylish and trendy

- ▶ all instructions are 32-bit
- ▶ 32 integer instruction including mret
- ▶ 6 instruction formats
- ▶ 6 immediate formats, sign extended, weirdly built, 12-bit or 20-bit
- ▶ all instructions update rd
- ▶ no delay slot
- ▶ branch target computed using pc as is
- ⇒ no specific implementation strategy inferred

Basic computer architecture



Basic computer architecture

Conclusion

Complexity is alike, but RISC-V has a few interesting features:

- ▶ all instructions update rd
- ▶ no delay slot
- ▶ branch computed using pc directly
- ▶ strange ways to compute the immediats
example: branch offset $\leftarrow (\text{IR}_{31}^{20} \parallel \text{IR}_7 \parallel \text{IR}_{30\ldots 25} \parallel \text{IR}_{11\ldots 8} \parallel 0)$

And is much more attractive to our students

google "RISC-V processor" : "Environ 9.050.000 résultats (0,60 secondes)"

google "mips processor" : "Environ 3.380.000 résultats (0,62 secondes)"

Small demo

Sneaking into some code

FPGA demo

Introduction to asm programming (3rd year, ≈250 students)

Objective of the class

- ▶ writing simple programs in asm
- ▶ understand variable classes: data, heap, stack
- ▶ systematically translate 'C' statements in asm
- ▶ ABI oriented towards function calling conventions

What did we do in the past?

- ▶ used an x86 subset
guaranteed headache, 13 different ABI running in the wild!
- ▶ moved to MIPS ISA, excluding unaligned word accesses, and MIPS o32 ABI

Introduction to asm programming (3rd year, ≈250 students)

What do we provide the students?

- ▶ Cross-development environment
- ▶ QEMU mimicking system of the “Basic Computer Architecture” class
- ▶ C source code of the exercises

What do we expect from the students?

- ▶ assembly code written as literally as possible
e.g. as generated by gcc -O0
- ▶ agile usage of the cross-dev environment
particularly gdb using remote connection on QEMU
- ▶ ability to call and be called from C functions
- ▶ understand low-level interrupt service routines

Introduction to asm programming (3rd year, ≈250 students)

MIPS I ABI in a nutshell

- ▶ all instructions but stores have their results in right-hand-side
- ▶ registers have hardware and software names
\$0, ..., \$31 vs \$zero, \$at, \$v0, ..., \$a0, ..., \$t0, ..., \$s0, ...
- ▶ register \$zero writable but always reads as zero
- ▶ 4 first arguments in \$a0, ..., \$a3 and return value in \$v0
- ▶ all jumps followed by a nop to avoid explaining the delay slot
- ▶ macros make use of the implicit register \$1 (\$at)

Introduction to asm programming (3rd year, ≈250 students)

- ▶ macro accept weirdly written asm lines

addi \$t0, 0xf00d ⇒ addi \$t0, \$t0, 0xf00d

lw and sw accept constants and symbols as arguments

lb \$t0, 0xdeadbeef ⇒ li \$t0, 0xdeadbeef / lb \$t0, 0(\$t0)

sw \$t0, variable ⇒ la \$at, variable / sw \$t0, 0(\$at)

But who knows if that is what the student meant to write ???

- ▶ sign extended or zero extended 16-bit immediates

visual instruction type and constant binary decoding easy for the teacher

addi \$v0,\$v0,0xffffdead ⇒ “operand out of range” although sign extended

ori \$v0,\$v0,0xdead ⇒ no error although zero extended

- ▶ interrupt/exception/traps all jump at same address

⇒ everything done in software

- ▶ access to cause, status and timer related registers using two simple instructions mfc0/mtc0

Introduction to asm programming (3rd year, ≈250 students)

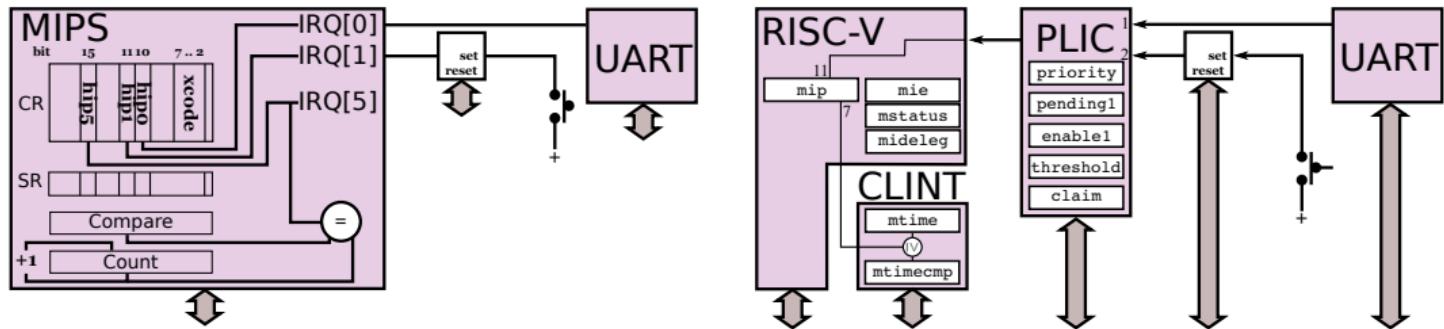
RISC-V psABI quite similar in spirit to MIPS

- ▶ all instructions but stores have their results in right-hand-side
- ▶ registers have hardware and software names
 x_0, \dots, x_{31} vs zero, ra, ..., a₀, ..., t₀, ..., s₀, ...
- ▶ register zero writable but always reads as zero
- ▶ 20 and 12 bit immediates always sign extend
addi x31, x31, 0xbad ⇒ "invalid operand error"
 $0x00000bad \neq 0xfffffbad$
- ▶ even stranger:

```
li x31, 0xdeadbeef => lui x31,      0xdeadc
                           addi x31, x31, 0xfffffeef
```
- ▶ sw and lw accept a symbol as argument, but not a constant

Introduction to asm programming (3rd year, ≈ 250 students)

- ▶ 8 first arguments in a_0, \dots, a_7 and return value in a_0
- ▶ no delay slot to hide
- ▶ access to cause, status and timer related registers using `csrrw` instructions
- ▶ only machine mode interrupt/exception/traps presented,
configuration without vectors \Rightarrow everything done in software
- ▶ CLINT and PLIC make things more complex than MIPS to handle simple timer interrupt



Introduction to asm programming (3rd year, ≈250 students)

Stack layouts

MIPS I r3000

high addresses ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ low addresses	5th parameter	other ones above needed only if callee wants them back after calling an other function \$sp in caller (f)
	room for \$a3	
	room for \$a2	
	room for \$a1	
	room for \$a0	
	\$ra	return address (g)
	registers to be saved	
	local variables	
	parameter n – 1	when calling an other function (h)
	...	
	parameter 5	
	room for \$a3	
	room for \$a2	
	room for \$a1	
	room for \$a0	for the next call (h) \$sp in callee (g)

RISC-V rv32i

high addresses ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ low addresses	f local variables	sp in caller (f) return address
	ra	
	other registers	
	to save	
	g local variables	
	parameter n – 1	preparing call to next function h
	...	
	parameter 8	sp in callee (g)
	...	

A bit easier to explain to students

Introduction to asm programming (3rd year, ≈250 students)

Conclusion

Complexity is alike, but RISC-V rv32i has a few interesting features:

- ▶ (very) low integer instruction count
- ▶ integer calling conventions pretty simple for fixed number of parameters
- ▶ no delay slot
- ▶ an unusual pc-relative instruction auipc

Small demo

Small function example

QEMU demo

Operating System Implementation (4th year, ≈ 75 students)

Objective of the class

- ▶ understand what happens when a computer is turned on
- ▶ understand virtual memory to physical memory translation
- ▶ learn to implement:
boot, kernel threads,
page tables, frame allocation (overlays \Rightarrow no file system, no page faults)
user processes, queues, shared memory, ...
- ▶ do all that on RISC-V 64 using QEMU sifive_u board

What did we do (and still do) before?

- ▶ originally developed for x86-32
- ▶ hiding quite a few stuff under the carpet
- ▶ still in use: only 2 groups over 8 did it on RISC-V last year

Operating System Implementation (4th year, ≈ 75 students)

What do we provide the students?

- ▶ skeleton of code, with set of (complex) Makefiles
- ▶ header files with helper functions
inline asm stuff, Linux priority queues, ...
- ▶ ≈ 15 userland tests stressing the implementation

What do we expect from the students?

- ▶ handle timer interrupts
- ▶ build all OS functions first in kernel mode
starting by process creation and context switch
- ▶ add (limited) support for virtual memory
- ▶ understand that there is one page table structure per process
- ▶ add system calls to wrap the kernel OS functions

Operating System Implementation (4th year, ≈ 75 students)

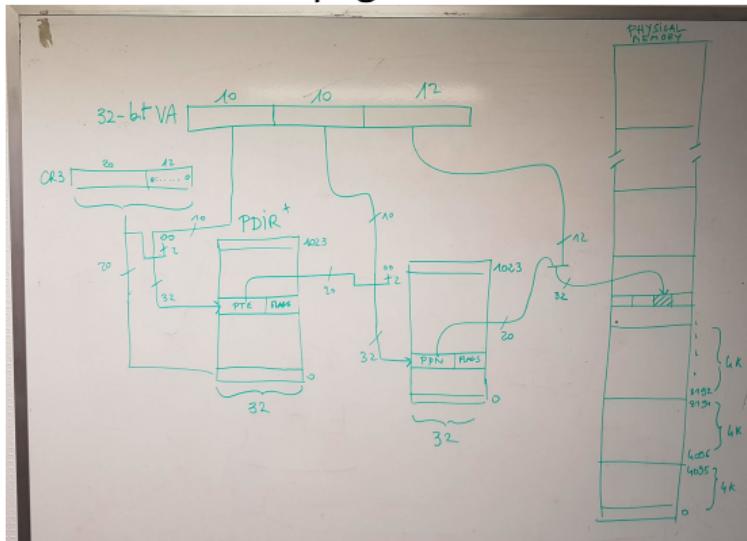
What changes when using RISC-V

- ▶ move to 64-bit
- ▶ boot phase can be fully written by the students
 - no gory details of x86 legacy to skip over
 - avoid hardwired stuffs that are hardly explainable: idt, gdt, tss, ...
 - ⇒ still need to configure the pmp areas
- ▶ simpler interruption mechanism and implementation of system calls
 - no implicit push of things on stack
 - ⇒ interruption setup a bit complex: m[ei]deleg, shadow registers, ...

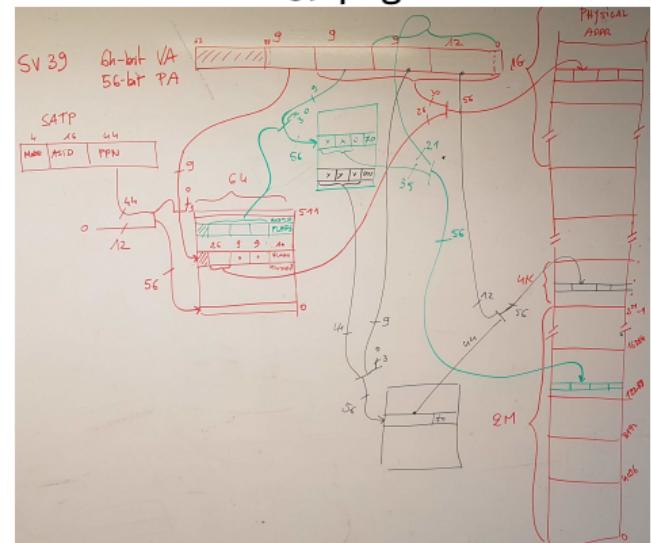
Operating System Implementation (4th year, ≈ 75 students)

And page tables, somehow more complex on RISC-V

x86 page tables



RISC-V SV39 page tables



Operating System Implementation (4th year, ≈ 75 students)

Conclusion

x86-32 vs RISC-V 64 fight not over!

- x86 students like x86 as they know its name

- x86 page table structure

- x86 see only kernel and user mode

- RISC-V boot process from start address user mode

- RISC-V no weird hw tables to update here and there

- RISC-V no implicit stuff pushed on or popped from the stack

Overall different, allowing to go a bit more in depth on hw related matters

Small demo

Setting and activating the page tables

Computer architecture (4th year, ≈ 35 students)

Objective of the class

- ▶ detail how simple 3/4/5/6 stage pipeline processors work
bypasses, interlocks
- ▶ explain cache coherency and memory consistency issues and solutions
MSI, MESI protocols, plus atomic operation support

What did we do (and still do) before?

- ▶ CAAQA with 5-stage pipeline MIPS, like everyone does!

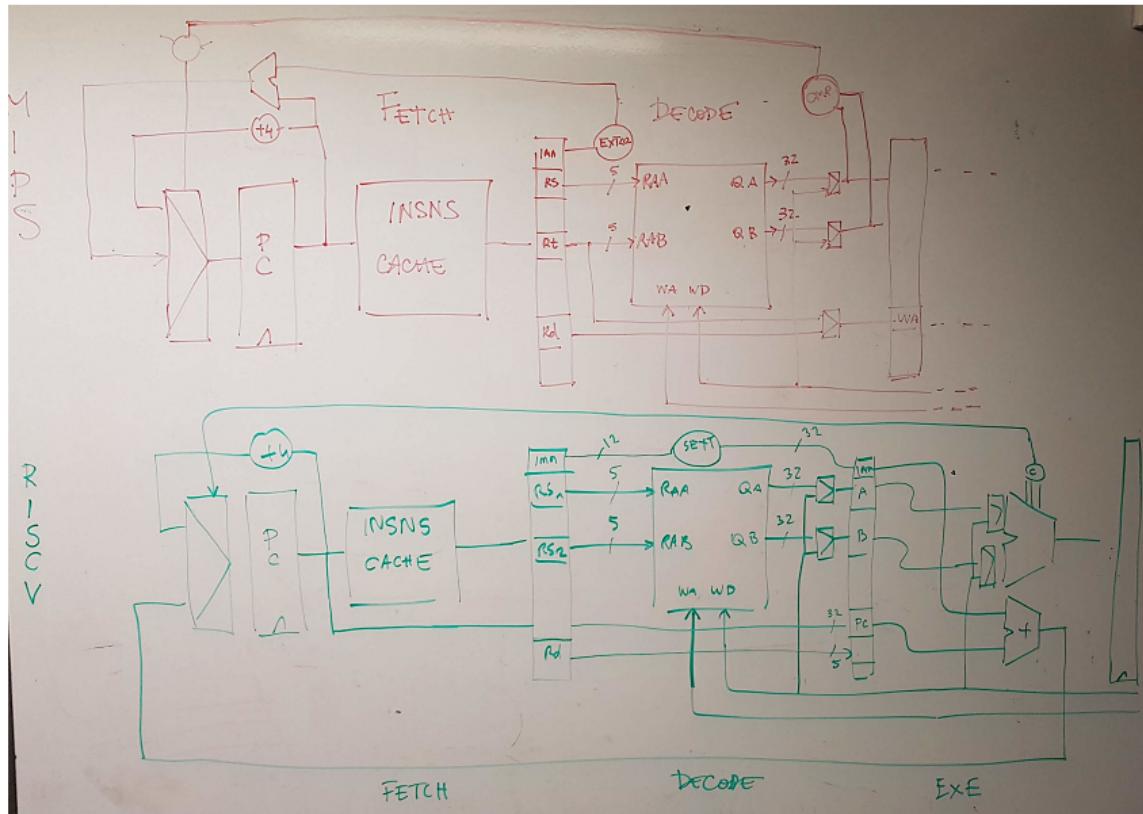
Computer architecture (4th year, ≈ 35 students)

Main changes due to RISC-V in classical 5-stage pipeline

On branches

- ▶ non-conditional branches have no delay slot
kill (at least) 1 following instruction
 - ▶ conditional branches are complex and resolved late
kill 2 following instructions if branch taken
- ⇒ Quite expensive, ...

Computer architecture (4th year, ≈ 35 students)



Computer architecture (4th year, ≈ 35 students)

Conclusion

Simple 5-stage pipeline implementation of RISC-V is:

- ▶ less efficient or harder
branch instructions will increase the CPI
or a branch predictor must be added
- ▶ a bit bigger, because pc must be propagated
- ▶ but the critical path should be shorter

Numerous atomic memory operations to present

No demo

Lucky you!

System level design in SystemC (5th year, ≈20 students)

Objective of the class

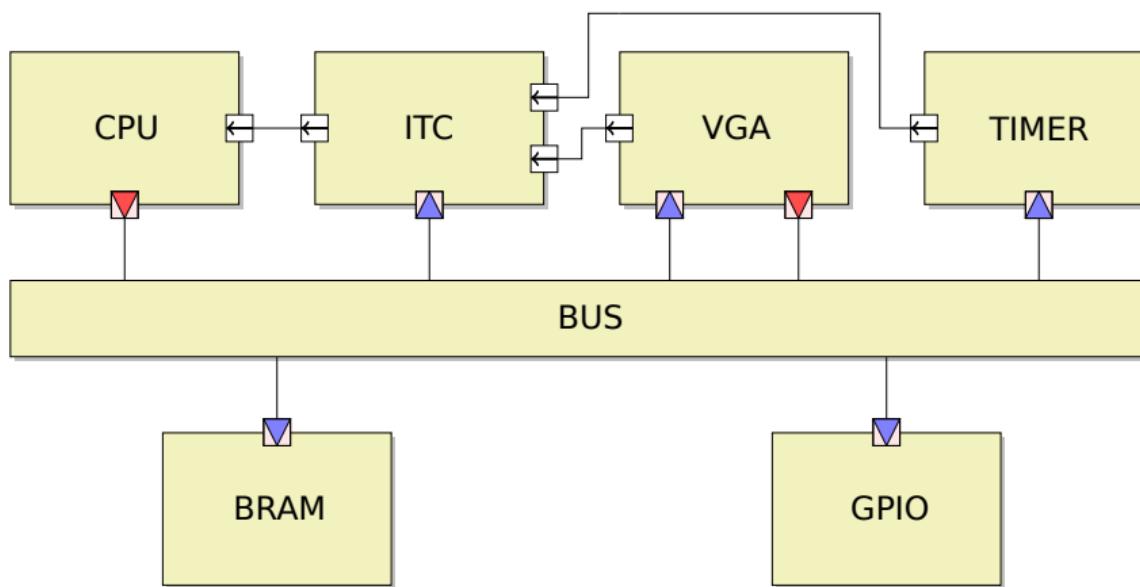
Class targeting an in-depth presentation of SystemC TLM concepts and modeling

- ▶ introduce the notion of virtual prototype
- ▶ show how to model and simulate a digital system in SystemC before its actual implementation
 - ▶ using "native" simulation, through dedicated APIs
 - ▶ using an instruction accurate simulator
- ▶ build the same system on a (cheap) FPGA board and run the same software

What did we do before?

- ▶ Originally developed around Xilinx' microblaze
 - ▶ SystemC model from the SoCLib library
 - ▶ microblaze RTL model from Xilinx

System level design in SystemC (5th year, ≈ 20 students)



System level design in SystemC (5th year, ≈ 20 students)

Using RISC-V

Interest

- ▶ very parameterizable architecture
- ▶ compressed instructions

But needs

- ▶ simple SystemC compatible ISS
⇒ developed a SoCLib rv32imafc model (machine mode only)
- ▶ synthesizable core on Xilinx FPGAs
⇒ have (a very very slow) one from Basic Architecture class, good enough

Note that keeping in pace with Vivado hurts, ...

Small demo

Native and cross-compiled top level

System execution

HW/SW System Integration with RISC-V (5th year, ≈40 students)

Objective of the class

Teaching the link between HW and SW in SoCs

- ▶ SW environment (assembly, linker, ISA simul., compil., debug)
- ▶ HW environment (SoC generation, emulation, ...)
- ▶ exceptions, interrupts and traps
- ▶ multi-tasking, multi-processing and memory coherence
- ▶ development and integration of a custom peripheral with its HAL
- ▶ HW and SW mapping ⇒ C application on FPGA

HW/SW System Integration with RISC-V (5th year, ≈ 40 students)

Needs

- ▶ tunable, fast to simulate HW/SW platform
- ▶ open source and synthesizable HW/SW SoC
⇒ FPGA evaluation
- ▶ platform widely accepted and well supported
⇒ Berkeley's RISC-V based Rocket Chip SoC generator

TLB: Translation Lookaside Buffer

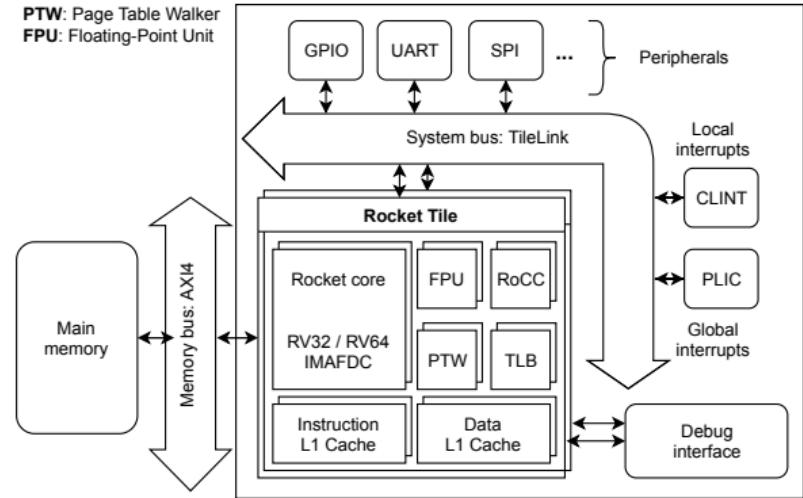
RoCC: Rocket Custom Coprocessor

PTW: Page Table Walker

FPU: Floating-Point Unit

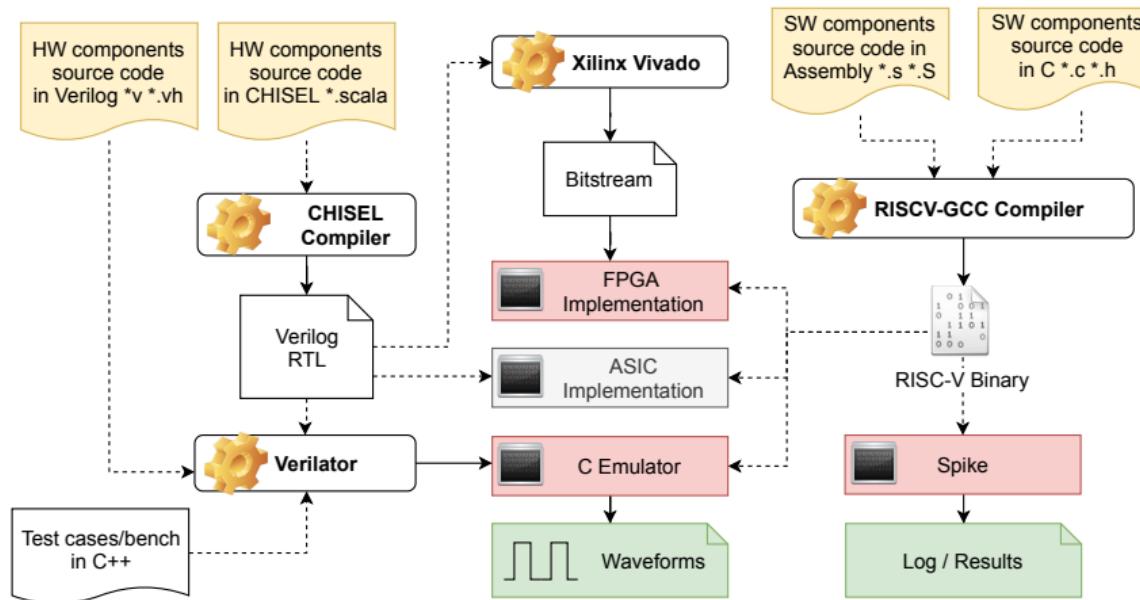
CLINT: Core Local Interruptor

PLIC: Platform-Level Interrupt Controller



HW/SW System Integration with RISC-V (5th year, ≈ 40 students)

Design Flow



HW/SW System Integration with RISC-V (5th year, ≈ 40 students)

Show impact of array placement in memory on L1 collision

L1 data-cache: 8 kB, 256 blocks, 32-byte each

x, y arrays of 4096 elements of 32-bit

address of x is 0x8101_0000, address of y is 0x8101_4000

```
loop: for (i = 0; i < 4096; i++)  
    s += x[i] + y[i];
```

- ▶ paper analysis of execution to evaluate dcache miss rate
- ▶ execution with spike to gather statistics
- ▶ propose a better placement of y to avoid collision
hint: make sure x and y never share a line

Small demo

Cache statistics with spike, with the bad and the good choices

HW/SW System Integration with RISC-V (5th year, ≈ 40 students)

Real-time multi-tasking

Students

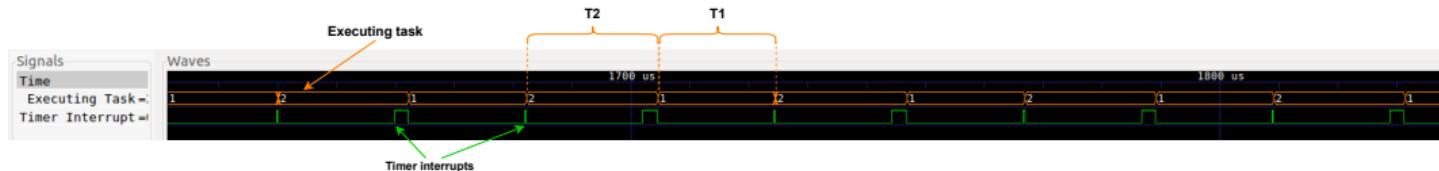
- ▶ analyze and complete assembly code
 - ▶ memory tasks allocation, init stacks, ...
 - ▶ save/restore contexts, switch tasks, scheduling, ...
 - ▶ time allocation: give #ticks per task
- ▶ code two simple tasks in C, e.g. increment in turn a shared variable
- ▶ simulate and debug using Spike simulator: functional validation

HW/SW System Integration with RISC-V (5th year, ≈ 40 students)

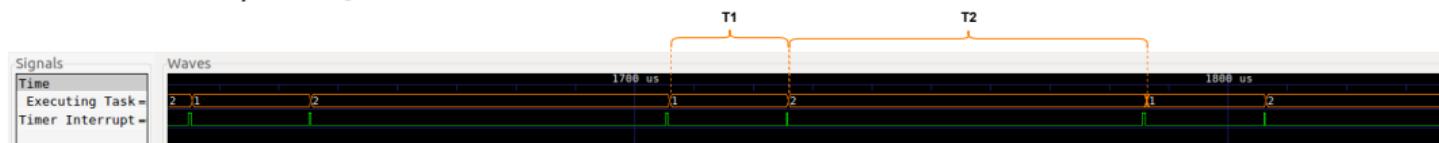
Real-time multi-tasking

- ▶ simulate the HW using the cycle-accurate bit-accurate C++ Emulator

$T_1 = T_2 = 100$ ticks



$T_1 = 100$ ticks, $T_2 = 300$ ticks



Outline

1 Preamble

2 Introduction

3 Overview of the curriculum

4 Class specifics

5 Take away

Take away

Momentum around RISC-V brings in some freshness

Opportunity to renew classes

Not a revolution: basics are basics, but:

- ▶ benefit from the hype
- ▶ make student aware that x86-64 doesn't rule them all
- ▶ escape complex CISC or complex RISC ISA

Benefit for teachers

- ▶ clean and orthogonal ISA (at least for rv32i)
- ▶ simple integer calling conventions
- ▶ stable cross-development tools, including simulators
- ▶ actual implementations on FPGA (even too many perhaps!) and ASIC

⇒ useful replacement to a mix of x86/MIPS/whatever all along the curriculum

Take away

Fear

Are we taking a reckless risk(-v)?

Thanks

Students/Engineers/PhD that have contributed to the classes

- ▶ Noureddine Ait-Said (IRT)
- ▶ Amaury Butaux (TIMA)
- ▶ Marius Leblanc (TIMA)
- ▶ Loïc Jovanovic (TIMA)
- ▶ Robin Stieglitz (IDEX)
- ▶ Arthur Vianes (IDEX)



Sponsors

